

Laplacian Pyramid for Image Encoding: Part 3

Finally, you are going to connect what you have implemented in the previous two labs with the code of this project. We will use the example image in the first step of the project and we will use the Laplacian encoding to compress it. We have set a predefined number of bins for each level of the pyramid, but you can select different values and see the effect in the image after the encoding.

For a comparison, we will also leverage code from the second lab, and approximate the image using Fourier transform. We will attempt to keep approximately the same compression rate for this technique, such that we need the same number of bits for each image (notice though that in both cases there are still options to optimize for the compression, although this is beyond the scope of this project).

Proceed with filling the last details of the code and visualizing the results using the script **Laplacian_encoding.m**. You can see the effect that different options have on the SNR of the image, as well as on quantitative aspects of the output (smoothing, high frequency details, noise, etc).

Your Script

 Save  Reset  MATLAB Documentation (<https://www.mathworks.com/help/>)

```
1 % load the image we will experiment with
2 I = imresize(double(rgb2gray(imread('lena.png'))),[256 256]);
3
4 % build the Laplacian pyramid of this image with 6 levels
5 depth = 6;
6 L = laplacianpyr(I,depth);
7
8 % compute the quantization of the Laplacian pyramid
9 bins = [16,32,64,128,128,256]; % number of bins for each pyramid level
10 LC = encoding(L,bins);
11
12 % compute the entropy for the given quantization of the pyramid
13 ent = pyramident(LC);
14
15 % Use the collapse command of the Lab 3 to recover the image
16 Ic = collapse(LC);
17
18 % compute the snr for the recovered image
19 snr_c = compute_snr(I,Ic);
20
21 % use the code from Lab 2 to compute an approximation image with
22 % the same level of compression approximately
23 [rows,cols] = size(I);
24 n_0 = rows*cols;
25 M = n_0/8;
26 Id = decompress(compress(I,sqrt(M)));
27 snr_d = compute_snr(I,Id);
28
29 % plot the resulting images
30 subplot(1,3,1);
31 imshow(I,[]); title('Original image');
32 subplot(1,3,2); imshow(Ic,[]);
33 title('Laplacian Encoding'); xlabel(['SNR = ' num2str(snr_c)]);
34 subplot(1,3,3); imshow(Id,[]);
35 title('Fourier Approximation'); xlabel(['SNR = ' num2str(snr_d)]);
36
37 function ent = pyramident(LC)
38
39     % Input:
40     % LC: the quantized version of the images stored in the Laplacian pyramid
41     % Output:
42     % br: the bitrate for the image given the quantization
43
44     % Please follow the instructions to fill in the missing commands.
45
46     ent = 0; % initialization of entropy
47     [r, c] = size(LC{1});
```

```

48     pixI = r*c;                % number of pixels in the original image
49
50     for i = 1:numel(LC)
51
52         % 1) Compute the number of pixels at this level of the pyramid
53         [r, c] = size(LC{i});
54         pix_i = r*c;
55
56         % 2) Compute the entropy at this level of the pyramid
57         % (MATLAB command: entropy)
58         ent_i = entropy(LC{i});
59
60         % 3) Each level contributes to the entropy of the pyramid by a
61         % factor that is equal to the sample density at this level, times
62         % the entropy at this level. The sample density is computed as
63         % (number of pixels at this level)/(number of pixels of original image).
64         % Add this to the current sum of the entropy 'ent'
65         ent = ent + (pix_i/pixI)*ent_i;
66
67     end
68
69 end
70
71 function LC = encoding(L, bins)
72
73     % Input:
74     % L: the Laplacian pyramid of the input image
75     % bins: [an array of ints, position i representing the]
76     % number of bins used for discretization of each pyramid level
77     % Output:
78     % LC: the quantized version of the image stored in the Laplacian pyramid
79
80     % Please follow the instructions to fill in the missing commands.
81
82     depth = numel(bins);
83     LC = cell(1,depth);
84
85     for i = 1:depth
86
87         % 1) Compute the edges of the bins we will use for discretization
88         % (MATLAB command: linspace)
89         % For level i, the linspace command will give you a row vector
90         % with bins(i) linearly spaced points between [X1,X2].
91         % Remember that the range [X1,X2] depends on the level of the
92         % pyramid. The difference images (levels 1 to depth-1) are in
93         % the range of [-128,128], while the blurred image is in the
94         % range of [0,256]
95         if i == depth % blurred image in range [0, 256]
96             edges = linspace(0,256,bins(i));
97         else % difference image in range [-128,128]
98             edges = linspace(-128,128,bins(i));
99         end
100
101         % 2) Compute the centers that correspond to the above edges
102         % The 1st center -> (1st edge + 2nd edge)/2
103         % The 2nd center -> (2nd edge + 3rd edge)/2 and so on
104         half_lng_int = (edges(2)-edges(1))/2;
105         centers = edges + half_lng_int .* ones(1,bins(i));
106         centers(end) = [];
107
108         % 3) Discretize the values of the image at this level of the
109         % pyramid according to edges (MATLAB command: discretize)
110         % Hint: use 'centers' as the third argument of the discretize
111         % command to get the value of each pixel instead of the bin index.
112
113         LC{i} = discretize(L{i}, edges, centers);
114
115     end

```

```

116     end
117
118 end
119
120 function I = collapse(L)
121
122     % Input:
123     % L: the Laplacian pyramid of an image
124     % Output:
125     % I: Recovered image from the Laplacian pyramid
126
127     % Please follow the instructions to fill in the missing commands.
128
129     depth = numel(L);
130
131     % 1) Recover the image that is encoded in the Laplacian pyramid
132     for i = depth:-1:1
133         if i == depth
134             % Initialization of I with the smallest scale of the pyramid
135             I = L{i};
136         else
137             % The updated image I is the sum of the current level of the
138             % pyramid, plus the expanded version of the current image I.
139             I = L{i} + expand(I);
140         end
141     end
142
143 end
144
145 function L = laplacianpyr(I,depth)
146
147     % Input:
148     % I: the input image
149     % depth: number of levels of the Laplacian pyramid
150     % Output:
151     % L: a cell containing all the levels of the Laplacian pyramid
152
153     % Please follow the instructions to fill in the missing commands.
154
155     L = cell(1,depth);
156
157     % 1) Create a Gaussian pyramid
158     % Use the function you already created.
159     G = gausspyr(I,depth);
160
161     % 2) Create a pyramid, where each level is the corresponding level of
162     % the Gaussian pyramid minus the expanded version of the next level of
163     % the Gaussian pyramid.
164     % Remember that the last level of the Laplacian pyramid is the same as
165     % the last level of the Gaussian pyramid.
166     for i = 1:depth
167         if i < depth
168             % same level of Gaussian pyramid minus the expanded version of next level
169             L{i} = G{i} - expand(G{i+1});
170         else
171             % same level of Gaussian pyramid
172             L{i} = G{i};
173         end
174     end
175
176 end
177
178 function G = gausspyr(I,depth)
179
180     % Input:
181     % I: the input image
182     % depth: number of levels of the Gaussian pyramid
183     % Output:

```

```

184 % G: a cell containing all the levels of the Gaussian pyramid
185
186 % Please follow the instructions to fill in the missing commands.
187
188 G = cell(1,depth);
189
190 % 1) Create a pyramid, where the first level is the original image
191 % and every subsequent level is the reduced version of the previous level
192 for i = 1:depth
193     if i == 1
194         G{i} = I; % original image
195     else
196         G{i} = reduce(G{i-1}); % reduced version of the previous level
197     end
198 end
199
200 end
201
202 function g = expand(I)
203
204     % Input:
205     % I: the input image
206     % Output:
207     % g: the image after the expand operation
208
209     % Please follow the instructions to fill in the missing commands.
210
211     % 1) Create the expanded image.
212     % The new image should be twice the size of the original image.
213     % So, for an n x n image you will create an empty 2n x 2n image
214     % Fill every second row and column with the rows and columns of the original image
215     % i.e., 1st row of I -> 1st row of expanded image
216     %         2nd row of I -> 3rd row of expanded image
217     %         3rd row of I -> 5th row of expanded image, and so on
218     I = im2double(I);
219     [m,n,clr] = size(I);
220     I_exp = zeros(2*m, 2*n, clr);
221     % note: 1:2 gives odd indices
222     I_exp(1:2:2*m, 1:2:2*n,:) = I(1:m, 1:n,:);
223
224     % 2) Create a Gaussian kernel of size 5x5 and
225     % standard deviation equal to 1 (MATLAB command fspecial)
226     Gauss = fspecial('gaussian',5,1);
227
228     % 3) Convolve the input image with the filter kernel (MATLAB command imfilter)
229     % Tip: Use the default settings of imfilter
230     % Remember to multiply the output of the filtering with a factor of 4
231     g = 4*imfilter(I_exp,Gauss);
232
233 end
234
235
236 function g = reduce(I)
237
238     % Input:
239     % I: the input image
240     % Output:
241     % g: the image after Gaussian blurring and subsampling
242
243     % Please follow the instructions to fill in the missing commands.
244
245     % 1) Create a Gaussian kernel of size 5x5 and
246     % standard deviation equal to 1 (MATLAB command fspecial)
247     Gauss = fspecial('gaussian',5,1);
248
249     % 2) Convolve the input image with the filter kernel (MATLAB command imfilter)
250     % Tip: Use the default settings of imfilter
251     I = im2double(I);

```

```

252     im_filtered = imfilter(I,Gauss);
253
254     % 3) Subsample the image by a factor of 2
255     % i.e., keep only 1st, 3rd, 5th, .. rows and columns
256     g = im_filtered(1:2:end, 1:2:end,:);
257
258 end
259
260 function [Fcomp] = compress(I,M_root)
261
262     % Input:
263     % I: the input image
264     % M_root: square root of the number of coefficients we will keep
265     % Output:
266     % Fcomp: the compressed version of the image
267
268     % Please follow the instructions in the comments to fill in the missing commands.
269
270     % 1) Perform the FFT transform on the image (MATLAB command fft2).
271     Fcomp = fft2(I);
272
273     % 2) Shift zero-frequency component to center of spectrum (MATLAB command fftshift).
274     Fcomp = fftshift(Fcomp);
275
276     % We create a mask that is the same size as the image. The mask is zero everywhere,
277     % except for a square with sides of length M_root centered at the center of the image.
278     [rows,cols] = size(I);
279     idx_rows = abs((1:rows) - ceil(rows/2)) < M_root/2 ;
280     idx_cols = abs((1:cols) - ceil(cols/2)) < M_root/2 ;
281     M = (double(idx_rows')) * (double(idx_cols));
282
283     % 3) Multiply in a pointwise manner the image with the mask.
284     Fcomp = Fcomp .* M;
285
286 end
287
288 function [Id] = decompress(Fcomp)
289
290     % Input:
291     % F: the compressed version of the image
292     % Output:
293     % Id: the approximated image
294
295     % Please follow the instructions in the comments to fill in the missing commands.
296
297     % 1) Apply the inverse FFT shift (MATLAB command ifftshift)
298     Id = ifftshift(Fcomp);
299
300     % 2) Compute the inverse FFT (MATLAB command ifft2)
301     Id = ifft2(Id);
302
303     % 3) Keep the real part of the previous output
304     Id = real(Id);
305
306 end
307
308 function snr = compute_snr(I, Id)
309
310     % Input:
311     % I: the original image
312     % Id: the approximated (noisy) image
313     % Output:
314     % snr: signal-to-noise ratio
315
316     % Please follow the instructions in the comments to fill in the missing commands.
317
318     % 1) Compute the noise image (original image minus the approximation)
319     noise_image = I - Id;

```

```
320
321 % 2) Compute the Frobenius norm of the noise image
322 noise_fnorm = norm(noise_image, 'fro');
323
324 % 3) Compute the Frobenius norm of the original image
325 orig_fnorm = norm(I, 'fro');
326
327 % 4) Compute SNR
328 snr = -20*log10(noise_fnorm/orig_fnorm);
329
330 end
331
```

[▶ Run Script](#)

Previous Assessment: All Tests Passed

[Submit](#)

✓ Is the estimated output for the image encoded with the Laplacian pyramid correct?

✓ Is the estimated output for the SNR of the Laplacian encoding correct?

✓ Is the estimated output for the image approximated with the 2D Fourier transform correct?

✓ Is the estimated output for the SNR of the Fourier approximation correct?

Output

